

```
# Dataset Link - https://www.kaggle.com/datasets/salader/dogs-vs-cats
```

```
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

▼ New Section

```
!kaggle datasets download -d salader/dogs-vs-cats
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.js
Downloading dogs-vs-cats.zip to /content
100% 1.06G/1.06G [00:10<00:00, 65.8MB/s]
100% 1.06G/1.06G [00:10<00:00, 107MB/s]
```

```
import zipfile
```

```
zip_data = zipfile.ZipFile('/content/dogs-vs-cats.zip')
zip_data.extractall('/content/')
zip_data.close()
```

Import Essential Libraries bold text

▼ New Section

```
import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

```
import matplotlib.pyplot as plt
import cv2
```

```
img = cv2.imread('/content/test/cats/cat.10057.jpg')
img
```

```
array([[136, 165, 169],
       [136, 165, 169],
       [136, 165, 169],
       ...,
       [251, 251, 251],
       [251, 251, 251],
       [251, 251, 251]],

       [[136, 165, 169],
       [136, 165, 169],
       [136, 165, 169],
       ...,
       [251, 251, 251],
       [251, 251, 251],
       [251, 251, 251]],

       [[135, 164, 168],
       [135, 164, 168],
       [135, 164, 168],
       ...,
       [251, 251, 251],
       [251, 251, 251],
       [251, 251, 251]],

       ...,

       [[103, 124, 126],
       [103, 124, 126],
       [103, 124, 126],
```

```

...,
[198, 226, 227],
[198, 226, 227],
[198, 226, 227]],


[[103, 124, 126],
[103, 124, 126],
[103, 124, 126],
...,
[198, 226, 227],
[198, 226, 227],
[198, 226, 227]],

[[104, 125, 127],
[104, 125, 127],
[104, 125, 127],
...,
[198, 226, 227],
[198, 226, 227],
[198, 226, 227]]], dtype=uint8)

```

```
plt.imshow(img)
```

```
<matplotlib.image.AxesImage at 0x7ce687f47820>
```



```
img.shape
```

```
(500, 331, 3)
```

```
img = cv2.imread('/content/test/cats/cat.10030.jpg')
img
```

```

array([[ 12,  10,  10],
       [ 14,  12,  12],
       [ 16,  14,  14],
       ...,
       [105, 118, 116],
       [ 85,  98,  96],
       [ 86,  99,  97]],

       [[ 11,   9,   9],
       [ 13,  11,  11],
       [ 14,  12,  12],
       ...,
       [101, 117, 116],
       [ 88, 102, 101],
       [ 91, 107, 106]],

       [[ 10,   8,   8],
       [ 11,   9,   9],
       [ 12,  10,  10],
       ...,
       [ 99, 120, 121],
       [ 91, 111, 112],
       [ 97, 118, 119]],

       ...,

       [[127, 135, 105],
       [128, 138, 108],
       [124, 136, 106],
       ...,
       [185, 161, 143],

```

```

[184, 160, 142],
[183, 159, 141]],


[[117, 127, 97],
[125, 137, 107],
[125, 140, 109],
...,
[182, 158, 140],
[181, 157, 139],
[180, 156, 138]],

[[112, 124, 94],
[125, 140, 109],
[130, 145, 114],
...,
[181, 157, 139],
[180, 156, 138],
[179, 155, 137]]], dtype=uint8)

```

```
plt.imshow(img)
```

```
<matplotlib.image.AxesImage at 0x7ce68c3a1540>
```



```
img.shape
```

```
(462, 500, 3)
```

```
# Generator
```

```
train_ds = tf.keras.utils.image_dataset_from_directory(
```

```

    directory = '/content/train',
    labels = 'inferred',
    label_mode = 'int',
    batch_size = 32,
    image_size = (256, 256)

```

```
)
```

```
test_ds = tf.keras.utils.image_dataset_from_directory(
```

```

    directory = '/content/test',
    labels = 'inferred',
    label_mode = 'int',
    batch_size = 32,
    image_size = (256, 256)

```

```
)
```

```

Found 20000 files belonging to 2 classes.
Found 5000 files belonging to 2 classes.

```

```
train_ds
```

```

<_PrefetchDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None,),
dtype=tf.int32, name=None))>

```

```
print(f'Number of Batches: {20000//32}')
```

```
Number of Batches: 625
```

```
0/255, 255/255
```

```
(0.0, 1.0)
```

```
# Normalization
```

```
def scale_down_px(image, label):
```

```
    image = tf.cast(image/255, tf.float32)
```

```
    return image, label
```

```
train_ds = train_ds.map(scale_down_px)
```

```
test_ds = test_ds.map(scale_down_px)
```

✓ Create a CNN model or Architecture

```
model = Sequential()
```

```
model.add(Conv2D(32, kernel_size=(3,3), padding='valid', activation='relu', input_shape= (256, 256, 3)))
```

```
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))
```

```
model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
```

```
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))
```

```
model.add(Conv2D(128, kernel_size=(3,3), padding='valid', activation='relu'))
```

```
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))
```

```
model.add(Flatten())
```

```
model.add(Dense(128, activation='relu'))
```

```
model.add(Dense(64, activation='relu'))
```

```
model.add(Dense(1, activation='sigmoid'))
```

```
model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d_3 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_4 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_4 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_5 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_5 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten_1 (Flatten)	(None, 115200)	0
dense_3 (Dense)	(None, 128)	14745728
dense_4 (Dense)	(None, 64)	8256
dense_5 (Dense)	(None, 1)	65

Total params: 14847297 (56.64 MB)
 Trainable params: 14847297 (56.64 MB)
 Non-trainable params: 0 (0.00 Byte)


```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')
```

```
history = model.fit(train_ds, validation_data=test_ds, epochs=10)
```

```
Epoch 1/10
625/625 [=====] - 72s 90ms/step - loss: 0.6436 - accuracy: 0.6248 - val_loss: 0.5800 - val_accuracy: 0.7094
Epoch 2/10
625/625 [=====] - 54s 86ms/step - loss: 0.5167 - accuracy: 0.7405 - val_loss: 0.5183 - val_accuracy: 0.7394
Epoch 3/10
625/625 [=====] - 55s 87ms/step - loss: 0.4081 - accuracy: 0.8081 - val_loss: 0.4481 - val_accuracy: 0.7948
Epoch 4/10
625/625 [=====] - 57s 91ms/step - loss: 0.2818 - accuracy: 0.8759 - val_loss: 0.5805 - val_accuracy: 0.7836
Epoch 5/10
625/625 [=====] - 59s 93ms/step - loss: 0.1562 - accuracy: 0.9387 - val_loss: 0.7151 - val_accuracy: 0.7976
Epoch 6/10
625/625 [=====] - 53s 84ms/step - loss: 0.0812 - accuracy: 0.9692 - val_loss: 1.0000 - val_accuracy: 0.7906
Epoch 7/10
625/625 [=====] - 53s 85ms/step - loss: 0.0665 - accuracy: 0.9762 - val_loss: 1.0412 - val_accuracy: 0.7892
Epoch 8/10
625/625 [=====] - 54s 86ms/step - loss: 0.0487 - accuracy: 0.9850 - val_loss: 1.2092 - val_accuracy: 0.7846
Epoch 9/10
625/625 [=====] - 54s 86ms/step - loss: 0.0376 - accuracy: 0.9876 - val_loss: 1.2022 - val_accuracy: 0.7846
Epoch 10/10
625/625 [=====] - 55s 88ms/step - loss: 0.0285 - accuracy: 0.9908 - val_loss: 1.3352 - val_accuracy: 0.7976
```


✓ Training/Validation Accuracy Graph

```
plt.plot(history.history['accuracy'], color='r', label='train')
plt.plot(history.history['val_accuracy'], color='b', label='validation')
plt.legend()
plt.show()
```



✓ Training/Validation loss Graph

```
plt.plot(history.history['loss'], color='red', label='train')
plt.plot(history.history['val_loss'], color='blue', label='validation')
plt.legend()
plt.show()
```



✓ Ways to improve model performance and to prevent overfitting

```
from keras.layers import BatchNormalization, Dropout

model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3), padding='valid', activation='relu', input_shape= (256, 256, 3)))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(128, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(1, activation='sigmoid'))

model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_6 (Conv2D)	(None, 254, 254, 32)	896
batch_normalization (Batch Normalization)	(None, 254, 254, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_7 (Conv2D)	(None, 125, 125, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 125, 125, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_8 (Conv2D)	(None, 60, 60, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 60, 60, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten_2 (Flatten)	(None, 115200)	0
dense_6 (Dense)	(None, 128)	14745728
dropout (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8256

```
dropout_1 (Dropout)      (None, 64)      0
dense_8 (Dense)          (None, 1)      65
```

```
=====
Total params: 14848193 (56.64 MB)
Trainable params: 14847745 (56.64 MB)
Non-trainable params: 448 (1.75 KB)
```


```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')
```

```
history = model.fit(train_ds, validation_data=test_ds, epochs=10)
```

```
Epoch 1/10
625/625 [=====] - 73s 111ms/step - loss: 1.4834 - accuracy: 0.5936 - val_loss: 0.5958 - val_accuracy: 0.6711
Epoch 2/10
625/625 [=====] - 69s 109ms/step - loss: 0.5925 - accuracy: 0.6874 - val_loss: 0.5440 - val_accuracy: 0.7111
Epoch 3/10
625/625 [=====] - 68s 108ms/step - loss: 0.5210 - accuracy: 0.7476 - val_loss: 0.6770 - val_accuracy: 0.5711
Epoch 4/10
625/625 [=====] - 67s 107ms/step - loss: 0.4510 - accuracy: 0.7893 - val_loss: 0.4255 - val_accuracy: 0.8011
Epoch 5/10
625/625 [=====] - 67s 107ms/step - loss: 0.3825 - accuracy: 0.8292 - val_loss: 0.4566 - val_accuracy: 0.7911
Epoch 6/10
625/625 [=====] - 68s 108ms/step - loss: 0.3237 - accuracy: 0.8623 - val_loss: 2.6355 - val_accuracy: 0.6611
Epoch 7/10
625/625 [=====] - 68s 108ms/step - loss: 0.2692 - accuracy: 0.8884 - val_loss: 0.6369 - val_accuracy: 0.7511
Epoch 8/10
625/625 [=====] - 68s 108ms/step - loss: 0.1857 - accuracy: 0.9269 - val_loss: 0.5086 - val_accuracy: 0.8111
Epoch 9/10
625/625 [=====] - 69s 110ms/step - loss: 0.1222 - accuracy: 0.9553 - val_loss: 0.6892 - val_accuracy: 0.8011
Epoch 10/10
625/625 [=====] - 67s 106ms/step - loss: 0.0864 - accuracy: 0.9679 - val_loss: 0.6227 - val_accuracy: 0.8211
```


Training/Validation Accuracy Graph

```
plt.plot(history.history['accuracy'], color='r', label='train')
plt.plot(history.history['val_accuracy'], color='b', label='validation')
plt.legend()
plt.show()
```



Training/Validation Loss Graph

```
plt.plot(history.history['loss'], color='red', label='train')
plt.plot(history.history['val_loss'], color='blue', label='validation')
plt.legend()
plt.show()
```

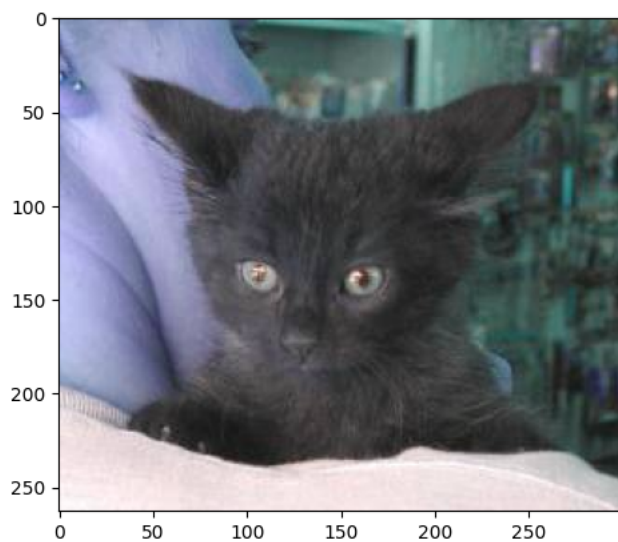


✓ Testing the Mdel

```
test_img = cv2.imread('/content/test/cats/cat.10036.jpg')
```

```
plt.imshow(test_img)
```

<matplotlib.image.AxesImage at 0x7ce68427b550>



```
test_img.shape
```

(263, 300, 3)

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```



```

<matplotlib.image.AxesImage at 0x7ce6842785b0>
0
50
test_img.shape
(256, 256, 3)
test_input = test_img.reshape(1, 256, 256, 3)
model.predict(test_input)
1/1 [=====] - 0s 246ms/step
array([[0.]], dtype=float32)
model.predict(test_input)[0]
1/1 [=====] - 0s 17ms/step
array([0.], dtype=float32)
model.predict(test_input)[0][0]
1/1 [=====] - 0s 20ms/step
0.0

output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')

if output >= 0.5:
    print('This is a Dog')
else:
    print('This is a Cat')

1/1 [=====] - 0s 27ms/step
Output is: 0.0

This is a Cat

```

```
test_img = cv2.imread('/content/test/dogs/dog.10032.jpg')
```

```
plt.imshow(test_img)
```

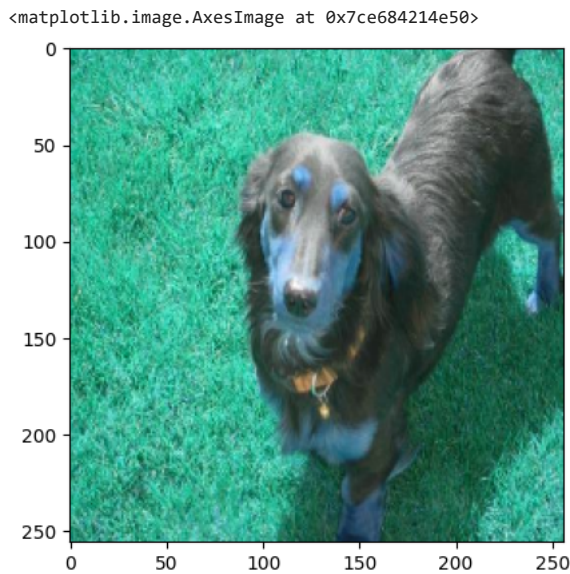


```
test_img.shape
```

```
(374, 500, 3)
```

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```



```
test_img.shape
```

```
(256, 256, 3)
```

```
test_input = test_img.reshape(1, 256, 256, 3)
```

```
model.predict(test_input)
```

```
1/1 [=====] - 0s 20ms/step
array([[1.]], dtype=float32)
```

```
model.predict(test_input)[0]
```

```
1/1 [=====] - 0s 19ms/step
array([1.], dtype=float32)
```

```
model.predict(test_input)[0][0]
```

```
1/1 [=====] - 0s 23ms/step
1.0
```

```
output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')
```

```
if output >= 0.5:
    print('This is a Dog')
else:
    print('This is a Cat')
```

```
1/1 [=====] - 0s 19ms/step
Output is: 1.0
```

```
This is a Dog
```

```
test_img = cv2.imread('/content/test/cats/cat.10243.jpg')
```

```
plt.imshow(test_img)
```

```
<matplotlib.image.AxesImage at 0x7ce684118700>
```




```
test_img.shape
```

```
(375, 499, 3)
```

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```

```
<matplotlib.image.AxesImage at 0x7ce68c2da8f0>
```



```
test_img.shape
```

```
(256, 256, 3)
```

```
test_input = test_img.reshape(1, 256, 256, 3)
```

```
model.predict(test_input)
```

```
1/1 [=====] - 0s 27ms/step  
array([[0.]], dtype=float32)
```

```
model.predict(test_input)[0]
```

```
1/1 [=====] - 0s 103ms/step  
array([0.], dtype=float32)
```

```
model.predict(test_input)[0][0]
```

```
1/1 [=====] - 0s 113ms/step  
0.0
```

```
output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')


if output >= 0.5:
    print('This is a Dog')
else:
    print('This is a Cat')

1/1 [=====] - 0s 33ms/step
Output is: 0.0

This is a Cat
```

```
test_img = cv2.imread('/content/test/dogs/dog.10865.jpg')
```

```
plt.imshow(test_img)
```



```
test_img.shape
```



```
(500, 442, 3)
```


```
test_img = cv2.resize(test_img, (256,256))
```


```
plt.imshow(test_img)
```

```

<matplotlib.image.AxesImage at 0x7ce686ac18a0>
test_img.shape

(256, 256, 3)
~ 
test_input = test_img.reshape(1, 256, 256, 3)

model.predict(test_input)

1/1 [=====] - 0s 38ms/step
array([[1.]], dtype=float32)

model.predict(test_input)[0]

1/1 [=====] - 0s 20ms/step
array([1.], dtype=float32)

model.predict(test_input)[0][0]

1/1 [=====] - 0s 22ms/step
1.0

output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')

if output >= 0.5:
    print('This is a Dog')
else:
    print('This is a Cat')

1/1 [=====] - 0s 20ms/step
Output is: 1.0

This is a Dog

```