

```
# Dataset Link - https://www.kaggle.com/datasets/biaiscience/dogs-vs-cats/download?
```

```
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

```
!kaggle datasets download -d biaiscience/dogs-vs-cats
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this,
Downloading dogs-vs-cats.zip to /content
100% 815M/817M [00:25<00:00, 39.3MB/s]
100% 817M/817M [00:25<00:00, 33.3MB/s]
```



```
import zipfile

zip_data = zipfile.ZipFile('/content/dogs-vs-cats.zip')
zip_data.extractall('/content/')
zip_data.close()
```

✓ Import Essential Libraries

```
import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

```
import matplotlib.pyplot as plt
import cv2
```

```
img = cv2.imread('/content/test/test/10039.jpg')
img
```

```
array([[112, 165, 168],
       [ 65, 115, 121],
       [ 87, 127, 139],
       ...,
       [ 73, 130, 121],
       [ 75, 132, 123],
       [ 82, 139, 130]],

      [[ 98, 151, 154],
       [ 56, 106, 112],
       [ 68, 110, 122],
       ...,
       [ 78, 135, 127],
       [ 77, 134, 126],
```

```

[ 80, 137, 129]],

[[ 92, 146, 147],
 [ 58, 109, 112],
 [ 60, 102, 114],
 ...,
 [ 75, 130, 127],
 [ 79, 134, 131],
 [ 86, 141, 138]],

...,

[[ 34, 60, 47],
 [ 21, 47, 33],
 [ 48, 74, 60],
 ...,
 [104, 159, 150],
 [ 81, 136, 127],
 [ 68, 123, 114]],

[[ 32, 56, 48],
 [ 15, 40, 30],
 [ 38, 62, 52],
 ...,
 [ 71, 126, 117],
 [ 55, 110, 101],
 [ 46, 101, 92]],

[[ 39, 63, 55],
 [ 18, 43, 33],
 [ 35, 59, 49],
 ...,
 [ 73, 128, 119],
 [ 64, 119, 110],
 [ 60, 115, 106]]], dtype=uint8)

```

```
plt.imshow(img)
```

```
<matplotlib.image.AxesImage at 0x7ac5a9bd5f30>
```



```
img.shape
```

```
(499, 388, 3)
```



```
# Generator
```

```
train_ds = tf.keras.utils.image_dataset_from_directory(
```

```
    directory = '/content/train',
    labels = 'inferred',
    label_mode = 'int',
    batch_size = 32,
    image_size = (256, 256)
```

```
)
```

```
test_ds = tf.keras.utils.image_dataset_from_directory(
```

```
    directory = '/content/test',
    labels = 'inferred',
    label_mode = 'int',
    batch_size = 32,
    image_size = (256, 256)
```

```
)
```

```
Found 25000 files belonging to 1 classes.
```

```
Found 12500 files belonging to 1 classes.
```

```
train_ds
```

```
<_PrefetchDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32, name=None))>
```

```
print(f'Number of Batches: {20000//32}')
```

```
Number of Batches: 625
```

```
0/255, 255/255
```

```
(0.0, 1.0)
```

```
# Normalization
```

```
def scale_down_px(image, label):

    image = tf.cast(image/255, tf.float32)

    return image, label
```

```
133/255
```

```
0.5215686274509804
```

```
train_ds = train_ds.map(scale_down_px)
test_ds = test_ds.map(scale_down_px)
```

✓ Create a CNN Model

```
model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3), padding='valid', activation='relu', input_shape=
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(128, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0

conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_2 (MaxPoolin g2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 128)	14745728
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 1)	65

```

=====
Total params: 14847297 (56.64 MB)
Trainable params: 14847297 (56.64 MB)
Non-trainable params: 0 (0.00 Byte)
=====

```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')
```

```
history = model.fit(train_ds, validation_data=test_ds, epochs=10)
```

```

Epoch 1/10
782/782 [=====] - 95s 105ms/step - loss: 9.0827e-04 - accuracy: 0.0000e+00
Epoch 2/10
782/782 [=====] - 80s 101ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 3/10
782/782 [=====] - 80s 102ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 4/10
782/782 [=====] - 99s 126ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 5/10
782/782 [=====] - 98s 125ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 6/10
782/782 [=====] - 98s 125ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 7/10
782/782 [=====] - 78s 100ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 8/10
782/782 [=====] - 99s 127ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 9/10
782/782 [=====] - 78s 99ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00
Epoch 10/10
782/782 [=====] - 77s 98ms/step - loss: 0.0000e+00 - accuracy: 0.0000e+00

```




✓ Training/Validation Accuracy Graph

```


plt.plot(history.history['accuracy'], color='r', label='train')
plt.plot(history.history['val_accuracy'], color='b', label='validation')
plt.legend()
plt.show()

```



✓ Training/Validation Loss Graph

```
plt.plot(history.history['loss'], color='red', label='train')
plt.plot(history.history['val_loss'], color='blue', label='validation')
plt.legend()
plt.show()
```



✓ Ways to Improve Model performance

```
from keras.layers import BatchNormalization, Dropout

model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3), padding='valid', activation='relu', input_shape=
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(128, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(1, activation='sigmoid'))

model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 254, 254, 32)	896
batch_normalization_3 (Batch Normalization)	(None, 254, 254, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_9 (Conv2D)	(None, 125, 125, 64)	18496
batch_normalization_4 (Batch Normalization)	(None, 125, 125, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_10 (Conv2D)	(None, 60, 60, 128)	73856
batch_normalization_5 (Batch Normalization)	(None, 60, 60, 128)	512

```
chNormalization)
```

```
max_pooling2d_8 (MaxPoolin (None, 30, 30, 128)      0
g2D)

flatten_2 (Flatten)      (None, 115200)      0

dense_6 (Dense)           (None, 128)      14745728

dropout_2 (Dropout)       (None, 128)      0

dense_7 (Dense)           (None, 64)      8256

dropout_3 (Dropout)       (None, 64)      0

dense_8 (Dense)           (None, 1)      65
```

```
=====
Total params: 14848193 (56.64 MB)
Trainable params: 14847745 (56.64 MB)
Non-trainable params: 448 (1.75 KB)
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')
```


```
history = model.fit(train_ds, validation_data=test_ds, epochs=10)
```

```
Epoch 1/10
782/782 [=====] - 95s 115ms/step - loss: 0.0041 - accuracy:
Epoch 2/10
782/782 [=====] - 113s 144ms/step - loss: 1.1480e-07 - accur
Epoch 3/10
782/782 [=====] - 91s 116ms/step - loss: 2.4390e-11 - accura
Epoch 4/10
782/782 [=====] - 91s 116ms/step - loss: 3.1752e-07 - accura
Epoch 5/10
782/782 [=====] - 91s 116ms/step - loss: 1.0840e-32 - accura
Epoch 6/10
782/782 [=====] - 89s 114ms/step - loss: 2.7913e-31 - accura
Epoch 7/10
782/782 [=====] - 89s 114ms/step - loss: 9.1050e-30 - accura
Epoch 8/10
782/782 [=====] - 91s 116ms/step - loss: 1.7425e-35 - accura
Epoch 9/10
782/782 [=====] - 90s 114ms/step - loss: 6.4295e-21 - accura
Epoch 10/10
782/782 [=====] - 92s 117ms/step - loss: 5.6011e-30 - accura
```



✓ Training/Validation Accuracy Graph


```
plt.plot(history.history['accuracy'], color='r', label='train')
plt.plot(history.history['val_accuracy'], color='b', label='validation')
plt.legend()
plt.show()
```



✓ Training/Validation Loss Graph

```
plt.plot(history.history['loss'], color='red', label='train')
plt.plot(history.history['val_loss'], color='blue', label='validation')
plt.legend()
plt.show()
```




✓ Testing the Model

```
test_img = cv2.imread('/content/test/test/10126.jpg')
```

```
plt.imshow(test_img)
```

<matplotlib.image.AxesImage at 0x7ac5236f5270>




```
test_img.shape
```

```
(332, 499, 3)
```

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```

```
<matplotlib.image.AxesImage at 0x7ac5247e2920>
```



```
test_img.shape
```

```
(256, 256, 3)
```

```
0 50 100 150 200 250
```

```
test_input = test_img.reshape(1, 256, 256, 3)
```

```
model.predict(test_input)
```

```
1/1 [=====] - 0s 387ms/step
array([[0.]], dtype=float32)
```

```
model.predict(test_input)[0]
```

```
1/1 [=====] - 0s 20ms/step
array([0.], dtype=float32)
```

```
model.predict(test_input)[0][0]
```

```
1/1 [=====] - 0s 18ms/step
0.0
```

```
output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')
```

```
if output >= 0.5:
    print('This is a Dog')
else:
    print('This is a Cat')
```


```
1/1 [=====] - 0s 27ms/step
Output is: 0.0
```

This is a Cat

```
test_img = cv2.imread('/content/test/test/10014.jpg')
```

```
plt.imshow(test_img)
```

<matplotlib.image.AxesImage at 0x7ac5241f3f70>




```
test_img.shape
```

```
(374, 500, 3)
```

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```

```
<matplotlib.image.AxesImage at 0x7ac5235253f0>
```



```
test_img.shape
```

```
(256, 256, 3)
```



```
test_input = test_img.reshape(1, 256, 256, 3)
```

```
model.predict(test_input)
```

```
1/1 [=====] - 0s 20ms/step
array([[0.]], dtype=float32)
```

```
model.predict(test_input)[0]
```

```
1/1 [=====] - 0s 21ms/step
array([0.], dtype=float32)
```

```
model.predict(test_input)[0][0]
```

```
1/1 [=====] - 0s 17ms/step
0.0
```

```
output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')
```

```
if output >= 0.5:
    print('This is a Dog')
else:
    print('This is a Cat')
```


```
1/1 [=====] - 0s 17ms/step
Output is: 0.0
```

```
This is a Cat
```

```
test_img = cv2.imread('/content/test/test/1000.jpg')
```

```
plt.imshow(test_img)
```

<matplotlib.image.AxesImage at 0x7ac5243d36d0>



```
test_img.shape
```

(499, 288, 3)

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```

```
<matplotlib.image.AxesImage at 0x7ac5243b9330>
```



```
test_img.shape
```

```
(256, 256, 3)
```



```
test_input = test_img.reshape(1, 256, 256, 3)
```



```
model.predict(test_input)
```

```
1/1 [=====] - 0s 20ms/step
array([[0.]], dtype=float32)
```

```
model.predict(test_input)[0]
```

```
1/1 [=====] - 0s 17ms/step
array([0.], dtype=float32)
```

```
model.predict(test_input)[0][0]
```

```
1/1 [=====] - 0s 18ms/step
0.0
```

```
output = model.predict(test_input)[0][0]
print(f'Output is: {output} \n')
```

```
if output >= 0.0:
    print('This is a Dog')
else:
    print('This is a Cat')
```


```
1/1 [=====] - 0s 19ms/step
Output is: 0.0
```

```
This is a Dog
```

```
test_img = cv2.imread('/content/test/test/10041.jpg')
```

```
plt.imshow(test_img)
```

<matplotlib.image.AxesImage at 0x7ac524382410>



```
test_img.shape
```

(437, 499, 3)

```
test_img = cv2.resize(test_img, (256,256))
```

```
plt.imshow(test_img)
```



```
<matplotlib.image.AxesImage at 0x7ac5241f3700>
```



```
test_img.shape
```

```
(256, 256, 3)
```



```
test_input = test_img.reshape(1, 256, 256, 3)
```



```
test_input = test_img.reshape(1, 256, 256, 3)
```



```
model.predict(test_input)
```

```
1/1 [=====] - 0s 19ms/step
array([[0.]], dtype=float32)
```

```
model.predict(test_input)[0]
```

```
1/1 [=====] - 0s 19ms/step
array([0.], dtype=float32)
```

```
model.predict(test_input)[0][0]
```

```
1/1 [=====] - 0s 19ms/step
0.0
```

```
output = model.predict(test_input)[0][0]
```

```
print(f'Output is: {output} \n')
```

```
if output >= 0.0:
```

```
    print('This is a Dog')
```

```
else:
```

```
    print('This is a Cat')
```

```
1/1 [=====] - 0s 17ms/step
Output is: 0.0
```

```
This is a Dog
```

